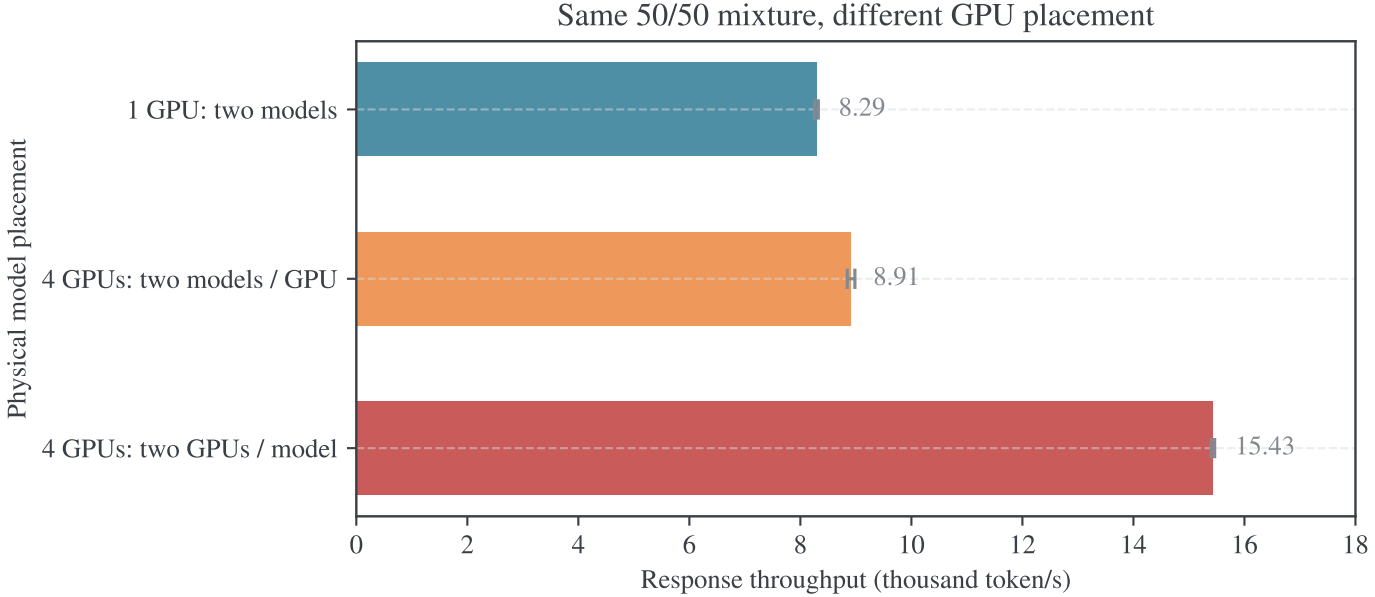
Same 50/50 mixture, different GPU placement
Physical model placement
1 GPU: two models
4 GPUs: two models / GPU
4 GPUs: two GPUs / model
8.29
8.91
15.43
0
2
4
6
8
10
12
14
16
18
Response throughput (thousand token/s)